

# logistic-regression

July 8, 2023

#INTERNSHIP PROJECT #TITLE : LOGISTIC REGRESSION

```
[1]: import pandas as pd
import numpy as np
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import \
    accuracy_score, confusion_matrix, classification_report
import matplotlib.pyplot as plt
import seaborn as sns
```

```
[2]: df= pd.read_csv('diabetes2.csv')
df.head() # Display the first few rows of the dataset
```

```
[2]:
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	\
0	6	148	72	35	0	33.6	
1	1	85	66	29	0	26.6	
2	8	183	64	0	0	23.3	
3	1	89	66	23	94	28.1	
4	0	137	40	35	168	43.1	

	DiabetesPedigreeFunction	Age	Outcome
0	0.627	50	1
1	0.351	31	0
2	0.672	32	1
3	0.167	21	0
4	2.288	33	1

```
[3]: df.info() # Get information about the dataset, including missing values
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 768 entries, 0 to 767
Data columns (total 9 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Pregnancies           768 non-null    int64
1   Glucose               768 non-null    int64
2   BloodPressure         768 non-null    int64
```

3	SkinThickness	768 non-null	int64
4	Insulin	768 non-null	int64
5	BMI	768 non-null	float64
6	DiabetesPedigreeFunction	768 non-null	float64
7	Age	768 non-null	int64
8	Outcome	768 non-null	int64


dtypes: float64(2), int64(7)  
memory usage: 54.1 KB

```
[12]: df.isnull().sum()
```

```
[12]: Pregnancies      0
      Glucose          0
      BloodPressure    0
      SkinThickness    0
      Insulin          0
      BMI              0
      DiabetesPedigreeFunction  0
      Age              0
      Outcome          0
      dtype: int64
```

```
[4]: sns.pairplot(df, hue="Outcome")
      plt.show
```

```
[4]: <function matplotlib.pyplot.show(close=None, block=None)>
```



```
[5]: x = df.drop("Outcome", axis=1)
y = df["Outcome"]
x_train, x_test, y_train, y_test = train_test_split(x, y, test_size=0.25,
↳random_state=42)
```

```
[6]: model = LogisticRegression()
model.fit(x_train, y_train)
y_pred = model.predict(x_test)
```

/usr/local/lib/python3.10/dist-packages/sklearn/linear\_model/\_logistic.py:458:  
ConvergenceWarning: lbfgs failed to converge (status=1):  
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max\_iter) or scale the data as shown in:  
<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:  
[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

```
n_iter_i = _check_optimize_result(
```

```
[7]: x_new=pd.DataFrame({'Pregnancies':[2], 'Glucose':[130], 'BloodPressure':  
    ↳ [60], 'SkinThickness':[35], 'Insulin':[95], 'BMI':  
    ↳ [35], 'DiabetesPedigreeFunction':[2.00], 'Age':[19]})
```

```
[8]: print(model.predict(x_new))
```

```
[1]
```

```
[9]: y_pred_train = model.predict(x)  
accuracy = accuracy_score(y, y_pred_train)  
print("Accuracy:", accuracy)
```

Accuracy: 0.765625

```
[10]: confusion_mat = confusion_matrix(y, y_pred_train)  
print("confusion matrix:")  
print(confusion_mat)
```

confusion matrix:

```
[[431  69]  
 [111 157]]
```

```
[11]: report = classification_report(y, y_pred_train)  
print("Classification Report:")  
print(report)
```

Classification Report:

	precision	recall	f1-score	support
0	0.80	0.86	0.83	500
1	0.69	0.59	0.64	268
accuracy			0.77	768
macro avg	0.74	0.72	0.73	768
weighted avg	0.76	0.77	0.76	768

```
[16]: plt.scatter(y_test,y_pred)  
plt.xlabel("Actual Outcome")  
plt.ylabel("predicted Outcome")  
plt.title("Actual vs Predicted Outcome")  
plt.show()
```

